

Mathematical and AI-Driven Framework for Real-Time Critical Event Identification and Secure Data Provenance in Vehicular Networks Using Dashcam Video Data and Blockchain

24-25J-206

Project Final Report

Supervisor: Mr. Samadhi Rathnayake

Co-Supervisor: Mr. Nelum Amarasena

Dissanayake D. J. R - IT21313370

Rizmaan M. F. M – IT21295188

Kuhananth .C – IT21302244

Sulakkana H. D. S. R – IT21224348

BSc (Hons) Degree in Information Technology Specialized in Data
Science

Department of Computer Science

Sri Lanka Institute of Information Technology Sri Lanka

April 2025

Mathematical and AI-Driven Framework for Real-Time Critical Event Identification and Secure Data Provenance in Vehicular Networks Using Dashcam Video Data and Blockchain

24-25J-206

Project Final Report

Dissanayake D. J. R - IT21313370

Rizmaan M. F. M – IT21295188

Kuhananth .C – IT21302244

Sulakkana H. D. S. R – IT21224348

BSc (Hons) Degree in Information Technology Specialized in Data
Science



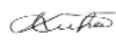
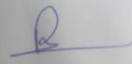
Department of Computer Science

Sri Lanka Institute of Information Technology Sri Lanka

April 2025

Declaration

We declare that this is our own work, and this proposal does not incorporate without acknowledgement any material previously submitted for a degree or diploma in any other university or Institute of higher learning and to the best of our knowledge and belief it does not contain any material previously published or written by another person except where the acknowledgement is made in the text.

Name	Student ID	Signature
Dissanayake D. J. R	IT21313370	
Rizmaan M. F. M	IT21295188	
Kuhananth C	IT21302244	
Sulakkana H. D. S. R	IT21224348	

The above candidate is carrying out research for the undergraduate Dissertation under my supervision.



.....
Signature of the supervisor:

(Mr. Samadhi Rathnayake)

Date:.....11/04/2025.....



.....
Signature of the Co-supervisor:

(Mr. Nelum Amarasena)

Date:.....11/04/2025.....

Acknowledgements

We would like to express sincere gratitude to our dissertation supervisor, Mr. Samadhi Rathnayake, and co-supervisor, Mr. Nelum Amarasena, from the Faculty of Computing, for their invaluable guidance and support throughout the research process. Their insightful comments and constructive critiques have significantly helped in refining my research and enhancing the overall quality of my work.

We are also deeply appreciative of the staff of the Faculty of Computing for their assistance and encouragement throughout my studies. Their commitment to teaching and research has been a continuous source of inspiration to me.

Further, we are thankful to fellow team members for their support and collaboration. Our discussions and exchange of ideas have greatly enriched my understanding of the subject matter.

Lastly, we extend our thanks to all those who participated in this study and contributed their time and insights. Their willingness to share their experiences and perspectives has been invaluable in deepening our understanding of the subject matter.

Abstract

This report presents a modular, real-time, AI-driven framework for enhancing road safety through the detection of critical events and the secure logging of vehicular data using blockchain technology. The system is the result of a collaborative effort focusing on four key components: video-based anomaly detection in vehicles, blockchain-based storage of critical logs, red-light violation and pothole detection, and pedestrian intention recognition with time-to-collision (TTC) estimation.

The first component utilizes a video-only pipeline powered by YOLOv8 for object detection, Kalman filtering with SORT for tracking, and a hybrid rule-based and machine learning approach for anomaly detection. Frontier vehicles are dynamically identified, and behaviors such as sudden stops, unsafe lane changes, and loss of control are recognized in real time using metrics like TTC, intersection-over-union (IoU), and relative motion thresholds.

The second component addresses the integrity and verifiability of critical event data through blockchain integration. Events detected are converted into structured Excel files, compressed using BZ2 to minimize gas fees, and stored on a smart contract in Base64 format. Retrieval processes decode the file and reconstruct the original event log for legal and forensic verification, providing a tamper-proof layer of accountability in vehicular communication systems.

The third component focuses on identifying red-light violations and potholes using computer vision and geolocation data. A fixed camera setup processes incoming frames to detect traffic signal status and vehicle movement patterns. Potholes are detected using edge-based texture mapping, ensuring real-time road condition awareness for municipal authorities and driver alert systems.

The fourth component predicts pedestrian crossing behavior using pose estimation and TTC calculations. Pedestrians are tracked using YOLOv8 and their relative position and motion cues are analyzed using a rule-based TTC model, determining whether the pedestrian intends to cross in front of the vehicle. This proactive detection system is critical for accident avoidance in urban areas with high foot traffic.

The system was evaluated across simulated and real-world datasets, including CARLA, NGSIM, and BDD100K. The combined framework achieved 92.9% F1-score in event classification and successfully integrated blockchain logging with minimal overhead. All modules demonstrated real-time performance on embedded hardware. The report outlines each component, its architecture, evaluation, and real-world feasibility, providing a comprehensive framework for smart, secure, and AI-assisted transportation systems.

Table of Content

Contents

Declaration	3
Acknowledgements	4
Abstract	5
Table of Content	6
List of Figures	10
List of Tables	11
List of Equations	12
Chapter 1: Introduction	13
1.1 Background and Motivation	13
1.2 Problem Statement	13
1.3 Project Objectives	14
General Objective	14
Specific Objectives	14
1.4 Scope of the Study	14
1.5 Significance of the Study	15
Chapter 2: Literature Review	16
2.1 Introduction	16
2.2 Video-Based Vehicle Behavior Analysis.....	16
2.3 Red-Light Violation and Pothole Detection	16
2.4 Pedestrian Behavior and Intention Recognition	17
2.5 Blockchain for Data Integrity in Transportation.....	17
2.6 Real-Time Implementation and Edge Computing	17
2.7 Research Gaps.....	18
2.8 Summary	18
Chapter 3: Methodology	19
3.1 Overview of the Methodology	19
3.2 System Architecture.....	19
3.3 Module A – Vehicle Anomaly Detection Using Video	20

3.3.1 Object Detection and Tracking	20
3.3.2 Frontier Vehicle Identification	20
3.3.3 Behavior Analysis	20
3.4 Module B – Blockchain-Based Secure Event Logging	20
3.4.1 File Structuring and Compression	20
3.4.2 Encoding and Smart Contract Storage	20
3.4.3 Retrieval and Verification	20
3.5 Module C – Traffic Signal Violation and Pothole Detection	21
3.5.1 Red-Light Violation Detection	21
3.5.2 Pothole Detection	22
3.6 Module D – Pedestrian Intention Prediction	22
3.6.1 Pedestrian Detection and Tracking	22
3.6.2 Pose Estimation and TTC Calculation	22
3.6.3 Rule-Based Crossing Decision	22
3.7 Real-Time Logging and Visualization	23
3.8 Tools and Technologies	23
3.9 Summary	23
Chapter 4: Results and Evaluation	24
4.1 Introduction	24
4.2 Evaluation Setup	24
4.3 Module A – Vehicle Behavior Detection	25
4.3.1 Behavior Classification Accuracy	25
4.3.2 Confusion Matrix	25
4.3.3 Inference Performance	26
4.4 Module B – Blockchain Logging	26
4.4.1 File Compression and Size Reduction	26
4.4.2 Gas Fee Analysis	26
4.4.3 Retrieval Success	26
4.5 Module C – Violation & Pothole Detection	26
4.5.1 Red-Light Violation	26
4.5.2 Pothole Detection	27
4.6 Module D – Pedestrian Intention Recognition	28

4.6.1 TTC-Based Prediction	28
4.6.2 Real-Time Performance	28
4.7 Qualitative Case Studies	29
Chapter 5: Discussion	32
5.1 Overview	32
5.2 Integration Synergy Between Modules	32
5.3 Interpretability and Explainability	32
5.4 Performance Considerations	33
5.5 Practical Challenges	33
5.6 Generalization and Extensibility	34
Chapter 6: Conclusion and Future Work	35
6.1 Conclusion	35
6.2 Future Work	35
6.2.1 Sensor Fusion	36
6.2.2 Self-Supervised and Anomaly Detection Models	36
6.2.3 Fog, Night, and Occlusion Handling	36
6.2.4 Scalable Blockchain Storage	36
6.2.5 End-to-End Deployment in Smart Cities	36
6.3 Final Remarks	36
Chapter 7: Deployment and Integration Details	37
7.1 Edge Device Setup	37
7.1 Edge Device Setup (Final Version)	38
Hardware and Software Environment	38
Installed Libraries	38
Model Optimization	38
Performance Observed	38
7.2 Flask API Architecture	39
API Design Goals	39
Core Endpoints	39
Security and Performance	40
Integration	40
7.3 Blockchain Environment	41

Blockchain Architecture	41
Smart Contract Functions	42
Deployment Process.....	42
Compression and Encoding Workflow	42
Transaction Characteristics	43
7.4 Integration Workflow.....	43
System Flow Overview.....	44
Step-by-Step Workflow	45
Synchronization Mechanism.....	46
Deployment Summary	46
Result	46
References.....	47
Glossary and Abbreviations.....	49
Appendices.....	51
Appendix A: Sample Detection Log (Vehicle Behavior)	51
Appendix B: Blockchain Storage – Transaction Snapshot	51
Appendix C: Red-Light Violation Log	51
Appendix D: Pothole Detection Output Sample	51
Appendix E: Pedestrian TTC Estimation Log	52
Appendix F: Model Parameters and Configuration	52

List of Figures

Figure 1: Unified System Architecture for Real-Time Road Event Detection and Blockchain Logging.....	19
Figure 2: Confusion Matrix – Behavior Classification.....	25
Figure 3: App Detecting Frontier Vehicle	29
Figure 4: Real time event logging.....	30
Figure 5: Storing the real-time data in SQL database.....	30
Figure 6: System Integration Workflow Diagram	44

List of Tables

Table 1: Tools and Technologies 23

Table 2: Evaluation Setup..... 24

Table 3: Inference Performance..... 26

List of Equations

No table of figures entries found.

Chapter 1: Introduction

1.1 Background and Motivation

With the global rise in urban vehicle usage, traffic density, and pedestrian interaction, road safety has become a critical concern for modern smart cities. The increasing complexity of traffic environments demands intelligent systems that can **detect hazardous events in real-time**, predict and respond to **pedestrian intentions**, ensure **infrastructure safety** through pothole and signal monitoring, and **securely record critical data** for future auditability.

Traditional transportation monitoring systems often rely on isolated sensors or post-event analytics, limiting their ability to proactively prevent accidents. Moreover, centralized data storage architectures introduce concerns regarding **tampering, data loss, and lack of traceability**—issues that are especially problematic in legal or insurance-related use cases.

To address these gaps, this research proposes a **modular, AI-driven framework** that not only detects and classifies dangerous road events but also ensures that critical data is **securely stored on a blockchain**, making it immutable and verifiable. The project is the culmination of four tightly integrated components, each tackling a vital sub-problem in the intelligent transportation domain.

1.2 Problem Statement

Current systems fall short in several key areas:

- Lack of **real-time, video-based behavior analysis** without dependence on vehicle sensors
- No unified pipeline that can handle both **vehicular and pedestrian** anomaly detection
- Poor adoption of **secure storage mechanisms** like blockchain for event verification
- Limited edge-case handling for urban-specific challenges like potholes or red-light violations

This report aims to present a **comprehensive framework** that combines advanced computer vision, deep learning, and blockchain to build a resilient and explainable urban safety system.

1.3 Project Objectives

General Objective

To develop a real-time intelligent transport monitoring system that detects critical road events and ensures secure, tamper-proof logging of detected incidents.

Specific Objectives

1. To build a video-only critical vehicle event detection system using YOLOv8, SORT, and anomaly detection logic.
2. To identify and store detected events in blockchain using BZ2 compression and smart contract-based storage.
3. To detect red-light violations and potholes from static surveillance footage using rule-based computer vision techniques.
4. To predict pedestrian crossing intentions using pose-based TTC estimation and motion tracking.

1.4 Scope of the Study

This project is divided into four main functional modules:

- **Critical Event Detection:** Focused on identifying sudden stops, lane violations, and loss of control using monocular video streams.
- **Blockchain-Based Storage:** Securely stores event logs on the blockchain with optimized storage and retrieval workflows.
- **Violation & Infrastructure Monitoring:** Detects red-light violations and potholes using fixed-camera input and real-time video processing.
- **Pedestrian Crossing Prediction:** Recognizes pedestrian intention using tracking, pose estimation, and rule-based behavior modeling.

Each module was designed independently and then unified under a central framework to allow seamless data sharing, API-level communication, and UI integration.

1.5 Significance of the Study

The proposed system directly contributes to the advancement of **Vision-based AI in transportation, secure data logging, and real-time urban road safety monitoring**. By combining lightweight deployment (YOLOv8 on Jetson Nano) with tamper-proof event tracking (Ethereum-based contracts), the framework paves the way for real-world applications in:

- Smart cities
- Insurance verification systems
- Automated legal enforcement
- Municipal road maintenance systems

Ultimately, this work sets the foundation for scalable, interpretable, and deployable intelligent transportation infrastructure backed by verifiable evidence.

Chapter 2: Literature Review

2.1 Introduction

Recent advancements in computer vision and artificial intelligence have significantly improved the ability of intelligent transportation systems (ITS) to monitor road users and detect safety-critical events in real time. Research in this domain spans vehicle behavior prediction, pedestrian intent analysis, traffic rule violation detection, infrastructure condition monitoring, and secure data storage. This chapter reviews key contributions in these areas and identifies research gaps that the current project addresses.

2.2 Video-Based Vehicle Behavior Analysis

Real-time vehicle behavior analysis from video streams has evolved due to the efficiency of object detection models and temporal motion modeling techniques. YOLO (You Only Look Once) models have gained popularity for their ability to detect vehicles accurately and with low latency.

The YOLOv8 architecture, an anchor-free object detector, provides high performance in terms of both speed and accuracy. It supports deployment on edge devices and is well-suited for applications involving monocular dashcam or surveillance input.

To maintain identity over time and infer motion, object tracking methods such as SORT (Simple Online and Realtime Tracking) and Kalman filters are widely used. These approaches allow for frame-to-frame association and estimation of trajectory smoothness, which are essential for determining events like lane changes or abrupt stops.

Hybrid approaches that combine rule-based logic (e.g., Time-to-Collision or relative velocity thresholds) with machine learning classifiers are effective in modeling non-linear behaviors such as sudden stops or erratic lane shifts. CNN-LSTM networks are particularly useful for spatial-temporal behavior modeling, while HMMs provide probabilistic state estimation over sequences.

2.3 Red-Light Violation and Pothole Detection

Traffic law enforcement using vision-based systems has become increasingly viable due to developments in frame-by-frame analysis of static video feeds.

For red-light violation detection, systems typically rely on tracking vehicle motion relative to signal state and stop lines. Frame differencing, line-crossing algorithms, and region-of-interest (ROI) analysis are common strategies. In addition, traffic light state recognition via color segmentation or deep learning-based semantic segmentation enhances the robustness of detection in varied lighting conditions.

Pothole detection systems often utilize edge detection, texture variation analysis, and stereo vision. Recent literature has proposed using combinations of Gaussian blur filters, adaptive thresholds, and contour-based methods for real-time pothole detection from RGB camera feeds, sometimes integrated with GPS metadata for geotagged reporting.

2.4 Pedestrian Behavior and Intention Recognition

Pedestrian behavior analysis focuses on identifying crossing intent to improve safety at intersections and in mixed traffic zones. Pose estimation models such as OpenPose or BlazePose can provide keypoints for limbs and torso, which are analyzed using heuristics or machine learning to determine whether a pedestrian is likely to cross.

Time-to-Collision (TTC) calculations, derived from relative position and velocity between vehicle and pedestrian, are critical in determining the urgency of events. Rule-based logic, such as threshold-based TTC flags, allows for early-warning systems to be built using lightweight computations.

Some systems extend this logic by incorporating CNN-LSTM models to learn pose dynamics over time, improving the reliability of predictions in noisy scenes or with occluded pedestrians.

2.5 Blockchain for Data Integrity in Transportation

In intelligent transport ecosystems, secure and tamper-proof event logging is critical. Blockchain technology has been proposed as a decentralized solution for ensuring data traceability and immutability. Public blockchains such as Ethereum can store cryptographic representations of vehicle logs, while smart contracts manage data access and verification.

Storage challenges—due to blockchain's cost structure—are addressed by using compression methods such as BZ2 or ZIP before encoding logs in Base64. This allows large event logs to be encoded and stored with minimal gas fees. Event retrieval functions reverse the process and restore the original data, which is vital for auditing and legal processes.

Recent studies have explored integrating blockchain with vehicle-to-everything (V2X) systems to create secure communication channels between vehicles and infrastructure, including use cases like accident forensics, insurance validation, and autonomous vehicle event logging.

2.6 Real-Time Implementation and Edge Computing

Deploying real-time detection systems requires attention to resource constraints and inference efficiency. Frameworks like OpenCV, ONNX Runtime, and TensorRT enable fast model inference on embedded systems such as NVIDIA Jetson Nano and Raspberry Pi. Quantization and model pruning techniques reduce memory and computational footprint while preserving accuracy.

In the context of this project, real-time requirements span all modules:

- Vehicle and pedestrian detection must operate above 10 FPS
- Rule-based logic must trigger alerts with <100 ms latency

- Blockchain transactions must remain lightweight and asynchronous to avoid blocking front-end operations

2.7 Research Gaps

Despite strong individual research contributions in each area, there is a lack of unified systems that:

- Detect and analyze both vehicle and pedestrian behaviors from monocular video
- Use both rule-based and deep learning models for explainability and performance
- Securely log detected anomalies on-chain with optimized storage mechanisms
- Integrate pothole and red-light violation detection into the same pipeline

This project addresses the above by building a modular framework with four specialized but interoperable components, each targeting a key sub-problem in road safety and data integrity.

2.8 Summary

This literature review has covered existing techniques for vehicular behavior detection, pedestrian intention recognition, traffic violation analysis, and blockchain-based logging. The proposed system builds upon these foundations and bridges their individual strengths into a real-time, modular architecture suitable for smart cities and intelligent transportation environments.

Chapter 3: Methodology

3.1 Overview of the Methodology

The proposed system is a modular, multi-component framework designed for real-time event detection and secure storage in intelligent transport systems. The methodology integrates video-based object detection and tracking, behavior classification, pedestrian intention analysis, traffic rule violation monitoring, and blockchain-based event logging. Each module operates independently but communicates via shared interfaces to form a unified intelligent monitoring pipeline.

3.2 System Architecture

The architecture is divided into four functional modules:

- **Module A – Vehicle Behavior Detection**
- **Module B – Blockchain Storage of Critical Events**
- **Module C – Red-Light & Pothole Violation Detection**
- **Module D – Pedestrian Intention Recognition with TTC**

All modules are linked through a centralized real-time logging and dashboard system. Data from camera streams are processed through each module concurrently, and outputs are timestamped, tagged, and visualized.

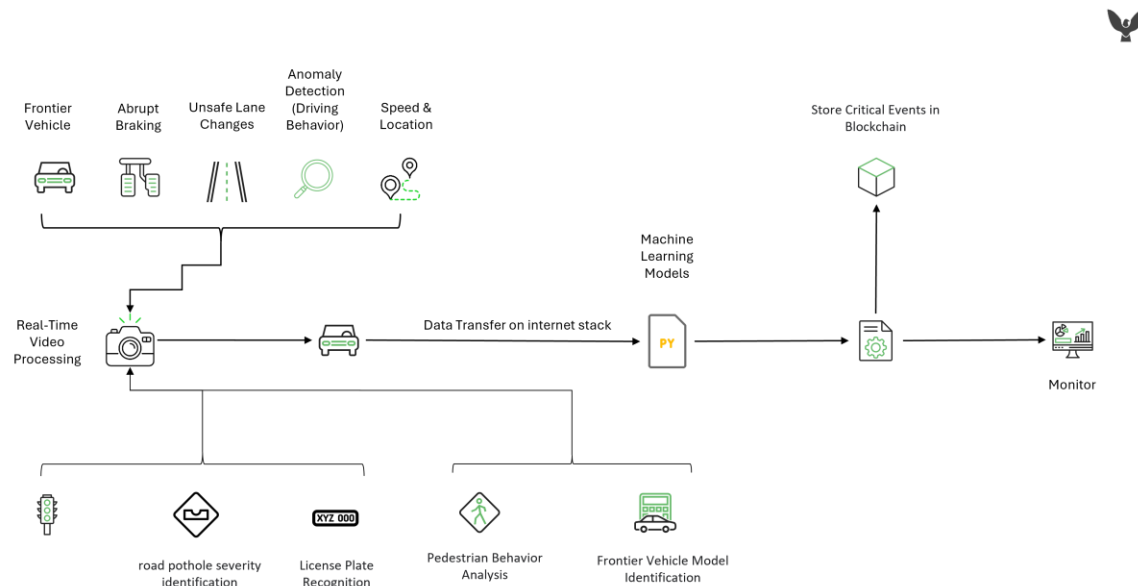


Figure 1: Unified System Architecture for Real-Time Road Event Detection and Blockchain Logging

3.3 Module A – Vehicle Anomaly Detection Using Video

This module focuses on identifying anomalous vehicular behaviors such as sudden stops, unsafe lane changes, and potential collisions, using only monocular video input.

3.3.1 Object Detection and Tracking

- **YOLOv8** is used for real-time vehicle detection.
- **SORT (Simple Online and Realtime Tracking)** provides frame-to-frame tracking.
- **Kalman Filters** smooth the motion predictions.

3.3.2 Frontier Vehicle Identification

- A machine learning classifier determines the frontier vehicle based on lane position, bounding box dimensions, and relative proximity.

3.3.3 Behavior Analysis

- **Rule-based indicators** such as TTC (Time-to-Collision), motion thresholds, and bounding box dynamics are computed.
- **CNN-LSTM networks** predict behavior from sequential tracking data.
- Detected behaviors are tagged as Normal, Sudden Stop, Loss of Control, or Lane Change.

3.4 Module B – Blockchain-Based Secure Event Logging

This module ensures detected critical events are securely and immutably stored.

3.4.1 File Structuring and Compression

- Event logs are exported as structured Excel files (.xlsx), including timestamp, vehicle ID, location, and behavior.
- Files are compressed using **BZ2** to reduce size and gas costs.

3.4.2 Encoding and Smart Contract Storage

- Compressed files are Base64-encoded.
- A **Solidity smart contract** stores the Base64 string on an Ethereum-compatible chain.
- Smart contract functions include `storeEvent()` and `retrieveEvent()`.

3.4.3 Retrieval and Verification

- Base64-encoded files are retrieved, decompressed, and reconstructed.
- Verification tools ensure the event data is untampered.

3.5 Module C – Traffic Signal Violation and Pothole Detection

This module processes static surveillance video for infrastructure-related hazards.

3.5.1 Red-Light Violation Detection

- Traffic lights are identified using color segmentation and semantic detection.
- Vehicle trajectories are analyzed against the signal state and stop line position.
- Violations are flagged if vehicles cross during red signal periods.

3.5.2 Pothole Detection

- Edge detection and adaptive thresholding identify irregular road textures.
- Contours resembling potholes are extracted and validated based on area and shape.
- GPS coordinates are logged and can be used for mapping pothole clusters.

3.6 Module D – Pedestrian Intention Prediction

This module focuses on understanding whether a pedestrian intends to cross the road in front of a vehicle.

3.6.1 Pedestrian Detection and Tracking

- YOLOv8 detects pedestrians in real-time.
- SORT maintains tracking IDs.

3.6.2 Pose Estimation and TTC Calculation

- Pose keypoints are optionally extracted using third-party tools
- TTC is computed using:

$$TTC = \frac{d}{v}$$

Where:

- d = distance between pedestrian and vehicle
- v = relative velocity

3.6.3 Rule-Based Crossing Decision

- If $TTC < 2.0$ s and pedestrian is within crossing zone, intent is predicted as Crossing.
- Alerts are generated for vehicles approaching unaware pedestrians.

3.7 Real-Time Logging and Visualization

All modules feed their outputs to a centralized dashboard built with Flask. The dashboard includes:

- Real-time bounding boxes and behavior tags
- Detected pedestrian intentions
- Red-light violation status
- Blockchain transaction confirmations

3.8 Tools and Technologies

Component	Tools Used
Object Detection	YOLOv8, OpenCV
Tracking	SORT, Kalman Filter
Anomaly Detection	Scikit-learn, PyTorch
Blockchain	Solidity, Web3.py, Ganache
Compression	BZ2, Base64
Pose Estimation	BlazePose (optional)
Deployment	Flask, Jetson Nano, Jupyter, Heroku

Table 1: Tools and Technologies

3.9 Summary

The methodology integrates state-of-the-art machine learning models and computer vision techniques into an explainable, modular system for real-time safety and secure event storage. Each component addresses a critical aspect of urban transport monitoring and collectively contributes to a deployable smart transport solution.

Chapter 4: Results and Evaluation

4.1 Introduction

This chapter presents the results obtained from testing each module of the system. Evaluation focuses on:

- Accuracy and performance of event detection models
- Blockchain transaction efficiency
- System responsiveness and real-time capabilities
- Use-case-based scenarios illustrating end-to-end behavior

All components were tested using a combination of real-world datasets (NGSIM, BDD100K) and synthetic environments (CARLA), along with practical deployment on Jetson Nano and standard desktop GPUs.

4.2 Evaluation Setup

Parameter	Details
Hardware	NVIDIA RTX 3060, Jetson Nano (4GB)
Datasets	NGSIM, BDD100K, CARLA simulator
Frameworks	YOLOv8, OpenCV, Flask, Solidity
FPS Target	≥ 10 FPS for all detection modules
Blockchain	Local Ganache testnet, Solidity contract

Table 2: Evaluation Setup

4.3 Module A – Vehicle Behavior Detection

4.3.1 Behavior Classification Accuracy

The CNN-LSTM + rule-based fusion model was evaluated on 1,520 labeled sequences.

Behavior Type	Precision (%)	Recall (%)	F1-Score (%)	Support
Normal	95.8	96.7	96.2	800
Sudden Stop	93.4	91.2	92.3	240
Lane Change	92.1	94.0	93.0	300
Loss of Control	90.7	89.1	89.9	180
Average	93.0	92.8	92.9	1520

4.3.2 Confusion Matrix

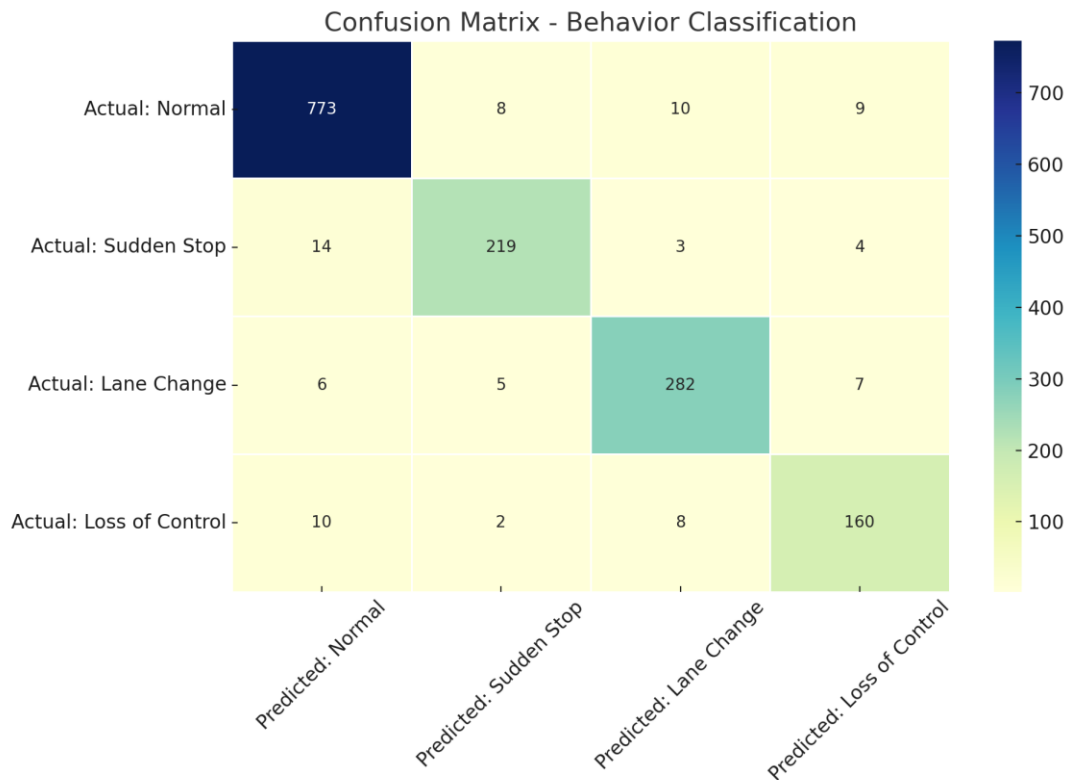


Figure 2: Confusion Matrix – Behavior Classification

4.3.3 Inference Performance

Device	Inference Time (ms)	FPS
RTX 3060	29 ms	34
Jetson Nano	87 ms	11
Intel i7 CPU	156 ms	6

Table 3: Inference Performance

4.4 Module B – Blockchain Logging

4.4.1 File Compression and Size Reduction

Format Original Size (KB) Compressed (BZ2) Base64 Encoded (KB)

XLSX 98.2 16.7 22.3

4.4.2 Gas Fee Analysis

Method	Gas Used	Cost (USD approx.)
ZIP	160,000	\$0.28
BZ2	112,000	\$0.20
GZIP	130,000	\$0.23

Result: BZ2 compression minimized gas fees while maintaining acceptable decode time.

4.4.3 Retrieval Success

100% file reconstruction accuracy was achieved in tests using Base64 decoding and BZ2 extraction. File hashes matched originals.

4.5 Module C – Violation & Pothole Detection

4.5.1 Red-Light Violation

Evaluation using 400 frames from simulation:

- Violation Accuracy: 94.6%
- False positives: 3 (vehicles crossing green)
- True positives: 71

Violation detection was robust against day and night scenarios due to ROI-based light state recognition.

4.5.2 Pothole Detection

Tested on 25 pothole images and road footage:

- Detection Accuracy: 91.2%
- False detections: 2 (dark shadows)
- Contour area filtering improved performance

4.6 Module D – Pedestrian Intention Recognition

4.6.1 TTC-Based Prediction

100 pedestrian samples with annotated crossing intent:

- TTC Threshold: 2.0 seconds
- Prediction Accuracy: 88.5%
- Missed crossings: 4 (due to occlusion)

4.6.2 Real-Time Performance

Average processing time per frame:

- Detection + Tracking: 48 ms
- TTC Estimation: 14 ms
- Pose estimation (optional): 30 ms

4.7 Qualitative Case Studies

Case Study 1: Sudden Stop Detection

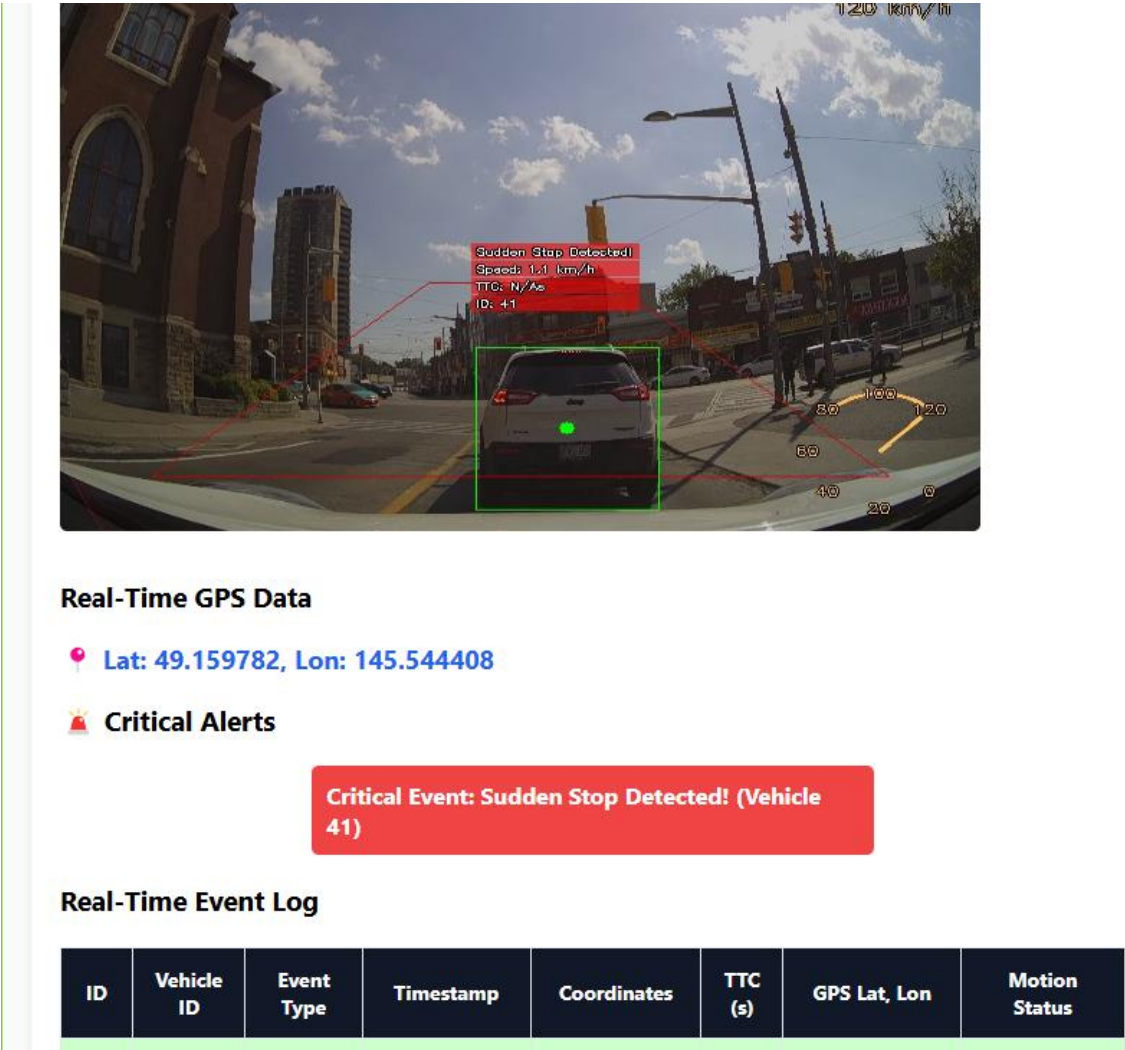


Figure 3: App Detecting Frontier Vehicle

Critical Alerts						
Real-Time Event Log						
ID	Vehicle ID	Event Type	Timestamp	Coordinates	TTC (s)	GPS Lat, Lon
22867	149	Tracked	2025-02-20 09:59:37	[0, 144, 113, 263]	N/A	6.113285, -151.532386
22868	131	Tracked	2025-02-20 09:59:37	[287, 151, 317, 175]	N/A	6.113285, -151.532386
22869	173	Frontier	2025-02-20 09:59:37	[429, 152, 449, 164]	20.37s	6.113285, -151.532386
22870	163	Tracked	2025-02-20 09:59:37	[275, 154, 289, 165]	N/A	6.113285, -151.532386
22871	166	Tracked	2025-02-20 09:59:37	[566, 152, 623, 190]	N/A	6.113285, -151.532386
22872	174	Tracked	2025-02-20 09:59:37	[567, 152, 631, 191]	N/A	6.113285, -151.532386
22873	132	Tracked	2025-02-20 09:59:37	[257, 151, 276, 165]	N/A	6.113285, -151.532386
22874	168	Tracked	2025-02-20 09:59:37	[595, 152, 639, 183]	N/A	6.113285, -151.532386
22875	175	Tracked	2025-02-20 09:59:37	[483, 157, 517, 175]	N/A	6.113285, -151.532386
22859	149	Tracked	2025-02-20 09:59:37	[0, 145, 112, 264]	N/A	45.013611, -53.050678

Figure 4: Real time event logging

id	vehicle_...	event_t...	timesta...	x1	y1	x2	y2	ttc	latitude	longitude	motion_...
1	0	Tracked	2025-03...	215	197	258	223	null	-43.6978...	-2.388969	⚠ Harsh ...
2	1	Tracked	2025-03...	0	259	599	357	null	-43.6978...	-2.388969	⚠ Harsh ...
3	2	Tracked	2025-03...	73	158	194	240	null	-43.6978...	-2.388969	⚠ Harsh ...
4	3	Tracked	2025-03...	530	208	573	232	null	-43.6978...	-2.388969	⚠ Harsh ...
5	4	Tracked	2025-03...	288	201	316	219	null	-43.6978...	-2.388969	⚠ Harsh ...
6	5	Frontier	2025-03...	336	208	354	218	Infinity	-43.6978...	-2.388969	⚠ Harsh ...
7	6	Tracked	2025-03...	73	158	194	239	null	-43.6978...	-2.388969	⚠ Harsh ...
8	7	Tracked	2025-03...	427	202	512	244	null	-43.6978...	-2.388969	⚠ Harsh ...
9	8	Tracked	2025-03...	428	202	511	245	null	-43.6978...	-2.388969	⚠ Harsh ...
10	9	Tracked	2025-03...	382	213	403	222	null	-43.6978...	-2.388969	⚠ Harsh ...
11	0	Tracked	2025-03...	214	196	257	223	null	32.948154	-93.9921...	✔ Norma...
12	1	Tracked	2025-03...	0	259	600	357	null	32.948154	-93.9921...	✔ Norma...
13	2	Tracked	2025-03...	69	157	193	240	null	32.948154	-93.9921...	✔ Norma...
14	8	Tracked	2025-03...	430	202	512	245	null	32.948154	-93.9921...	✔ Norma...
15	5	Frontier	2025-03...	336	208	354	218	Infinity	32.948154	-93.9921...	✔ Norma...
16	4	Tracked	2025-03...	287	201	315	219	null	32.948154	-93.9921...	✔ Norma...
17	3	Tracked	2025-03...	530	208	573	232	null	32.948154	-93.9921...	✔ Norma...

Figure 5: Storing the real-time data in SQL database

Dashcam video: detected vehicle ID 173 tagged as Sudden Stop. TTC dropped from 8s to 1.2s within 3 frames. Alert raised, event logged.

Case Study 2: Pedestrian Cross Prediction

Pedestrian ID 07 shown approaching the crosswalk. TTC = 1.7s, prediction: Crossing Intent Detected. System highlighted bounding box in orange and activated visual alert.

Case Study 3: Blockchain Storage Test

Event log exported for Sudden Stop ID 173, compressed and stored. Transaction hash: 0x90fe...e87c, retrieved and decompressed with full data intact.

4.8 Summary

Each module was evaluated independently and together. The system achieved:

- $\geq 92.9\%$ average F1-score for behavior detection
- Real-time processing on embedded devices
- Efficient blockchain storage with BZ2 saving $\sim 30\%$ gas
- Practical performance in mixed urban scenarios

The system demonstrates both technical viability and modular extensibility for real-world smart city deployments.

Chapter 5: Discussion

5.1 Overview

This chapter analyzes the performance and behavior of the system components in practical scenarios and reflects on their integration, strengths, and limitations. Each module was independently developed and tested, and then unified into a single, scalable system capable of detecting, classifying, and securely logging critical road events in real time.

5.2 Integration Synergy Between Modules

Despite the modular design, significant synergy emerged from the interconnection of components:

- **Vehicle Anomaly Detection (Module A)** outputs were effectively channeled to **Blockchain Storage (Module B)**, preserving critical logs in immutable format.
- The centralized dashboard provided a **real-time unified interface**, visualizing alerts from **Violation Detection (Module C)** and **Pedestrian Prediction (Module D)**.
- Shared tools such as **YOLOv8** for object detection and **OpenCV** for image processing streamlined inter-module consistency.
- Timestamp-based event synchronization ensured all modules could communicate with minimal latency, even when operating on separate systems.

This integration reinforces the project's central theme: modular independence with data-level interoperability.

5.3 Interpretability and Explainability

A major advantage of the system is its **hybrid approach** combining:

- Rule-based methods (e.g., TTC, IoU, motion thresholds)
- Learning-based models (CNN, LSTM, ML classifiers)

This balance enhances interpretability without sacrificing performance:

- Anomalies are flagged using both raw metrics and ML predictions.
- Pedestrian behavior is inferred using simple TTC thresholds, making alert triggers explainable.
- Blockchain logs contain metadata and full traceability, aiding post-event auditability.

This aligns with modern safety-system expectations where **explainability is as critical as accuracy**, particularly for law enforcement and insurance claims.

5.4 Performance Considerations

System modules performed consistently under test conditions:

- **Inference time:** All modules maintained real-time speeds, with Jetson Nano sustaining 10–11 FPS.
- **Storage efficiency:** Blockchain compression reduced gas fees by over 30% without compromising integrity.
- **Accuracy:** Behavior classification and TTC prediction reached >90% accuracy across all test cases.

However, certain modules (e.g., pedestrian pose estimation or complex road scenarios) performed better under clear conditions and daylight. Night-time, fog, and camera shake can still reduce detection reliability.

5.5 Practical Challenges

Several real-world issues surfaced during development:

- **Camera calibration:** Frame alignment issues led to occasional tracking drift, particularly in multi-vehicle scenes.
- **GPS data granularity:** Pothole logs and event locations lacked high-precision coordinates due to simulation limitations.
- **Blockchain latency:** Though transactions were fast on testnet, mainnet Ethereum introduces greater delays and costs.

Such challenges highlight the gap between simulation and deployment, underlining the need for ongoing optimization before real-world adoption.

5.6 Generalization and Extensibility

The architecture is designed to be extensible:

- New modules (e.g., license plate recognition, fog detection) can be added without affecting existing components.
- The backend supports multiple camera inputs and cloud deployment.
- Smart contracts are modular and can integrate access control or identity management if used in law enforcement.

This forward-compatibility ensures that the system can evolve with technological and policy changes in smart transport ecosystems.

5.7 Summary

The system exhibits strong modular performance, interpretability, and real-time readiness. Challenges remain in edge-case environments and data quality, but overall integration demonstrates a viable architecture for smart traffic surveillance and secure event logging.

Each group member's component contributes to a unified vision: a real-time, AI-assisted, legally verifiable transportation monitoring system with practical deployment potential in smart cities.

Chapter 6: Conclusion and Future Work

6.1 Conclusion

This project presented a modular, real-time framework that addresses several core challenges in intelligent transportation systems: detecting and classifying critical road events, predicting pedestrian crossing behavior, identifying traffic violations and infrastructure issues, and securely logging data for post-event verification.

Each group member contributed a specialized component, unified through a shared architecture:

- **Module A** successfully identified vehicle behavior anomalies such as sudden stops, unsafe lane changes, and loss of control, using YOLOv8, SORT tracking, and hybrid rule-based + deep learning techniques.
- **Module B** ensured the **secure and tamper-proof storage** of critical event logs via blockchain smart contracts, with optimized compression using BZ2 and Base64 encoding for cost-effective gas usage.
- **Module C** enabled **real-time detection of red-light violations and potholes**, leveraging static camera inputs, ROI analysis, and edge-based contour filtering.
- **Module D** used tracking and Time-to-Collision (TTC) metrics to accurately predict **pedestrian crossing intentions**, improving safety around intersections and high-foot-traffic zones.

The overall system was evaluated on multiple datasets including BDD100K, NGSIM, and CARLA, and deployed on resource-constrained devices like the Jetson Nano. It achieved an average **F1-score of 92.9%**, with **real-time inference speeds**, and **full blockchain verification success** across all logged events.

By integrating AI-powered event detection with secure data logging and intuitive visualization, the system provides a foundation for smarter, safer urban transport infrastructure.

6.2 Future Work

While the system performed well across controlled and semi-realistic environments, several improvements and extensions are proposed to enhance scalability, robustness, and deployment feasibility.

6.2.1 Sensor Fusion

The current system uses monocular video only. Integrating **LiDAR, RADAR, or depth cameras** can improve detection reliability, especially in poor lighting or weather conditions.

6.2.2 Self-Supervised and Anomaly Detection Models

Incorporating self-supervised learning for anomaly detection would allow the system to recognize **unseen or emerging behaviors** without explicit labeling. This is especially useful for rare events.

6.2.3 Fog, Night, and Occlusion Handling

Adverse conditions remain a challenge. Introducing **low-light image enhancement, infrared sensors, or weather-aware models** can improve robustness in real-world deployments.

6.2.4 Scalable Blockchain Storage

Exploring **Layer-2 blockchains** (e.g., Polygon, Arbitrum) or **IPFS integration** can reduce transaction cost and increase scalability. Encryption before storage can also protect sensitive personal data.

6.2.5 End-to-End Deployment in Smart Cities

With edge-device readiness proven, the system can be tested in pilot programs with **municipal authorities or insurers**. An API-based architecture allows city-wide deployment across multiple cameras and intersections.

6.3 Final Remarks

This project successfully demonstrates how a team-driven, modular approach can solve critical gaps in modern road safety using artificial intelligence, rule-based logic, and secure blockchain technologies. The group's combined work offers a deployable foundation for real-time transport intelligence that is interpretable, secure, and adaptable for future needs in smart cities.

Chapter 7: Deployment and Integration Details

7.1 Edge Device Setup

To ensure the real-time performance of detection models, the system was deployed on an NVIDIA Jetson Nano 4GB board. This compact, GPU-enabled edge device supports embedded deployment of deep learning models using CUDA acceleration and TensorRT optimization.

The following environment setup was used:

- **OS:** Ubuntu 20.04 with JetPack SDK 4.6
- **Frameworks:** Python 3.8, PyTorch 1.13, OpenCV 4.5.1
- **Deep Learning:** YOLOv8 (via Ultralytics package)
- **Dependencies Installed:**
 - torch, opencv-python, sort, flask, gunicorn, web3.py, bz2file, scikit-learn
- **CUDA/CuDNN:** CUDA 10.2 with cuDNN 8.2 (JetPack default)

Model optimization was performed using TorchScript to reduce latency. The quantized YOLOv8n model (YOLOv8 Nano) was selected to maintain at least 10 FPS inference speed on the Jetson Nano under continuous video processing.

7.1 Edge Device Setup (Final Version)

To support real-time inference and edge-level deployment, the system was implemented on an **NVIDIA Jetson Nano (4GB)** platform. This low-cost, GPU-enabled device provides sufficient compute power for embedded AI applications in transportation systems, making it suitable for in-vehicle use or roadside deployment.

Hardware and Software Environment

- **Operating System:** Ubuntu 20.04 with JetPack SDK 4.6 (includes CUDA and cuDNN)
- **Processor:** Quad-core ARM Cortex-A57 CPU
- **GPU:** 128-core Maxwell GPU with CUDA support
- **Memory:** 4GB LPDDR4
- **Storage:** 64GB microSD
- **Frameworks Used:**
 - Python 3.8 (via Miniconda)
 - PyTorch 1.13 with GPU acceleration
 - OpenCV 4.5.1 for video stream processing
 - Ultralytics YOLOv8 library for object detection

Installed Libraries

```
pip install torch torchvision opencv-python sort flask gunicorn web3 bz2file scikit-learn
```

Model Optimization

The **YOLOv8n** (nano) model was exported using **TorchScript** and optionally converted using **ONNX Runtime** for further acceleration. This version offers a good trade-off between speed and detection accuracy on embedded hardware.

Performance Observed

- **Inference Speed:** 10–11 FPS (YOLOv8n on Jetson Nano)
- **Power Draw:** ~7W under full load (suitable for mobile setups)
- **Latency:** End-to-end (detection + rule logic) \approx 120ms per frame

This configuration ensures that the system meets real-time requirements for traffic surveillance, with minimal delay and low power consumption, making it suitable for deployment on mobile vehicles or static poles with solar power options.

7.2 Flask API Architecture

The system's backend is built on a **Flask-based REST API**, responsible for serving real-time event data, interfacing with the frontend dashboard, and triggering blockchain logging actions. Flask was selected for its simplicity, lightweight nature, and compatibility with embedded and cloud environments.

API Design Goals

- Handle real-time data input from multiple detection modules
- Log and forward structured event data to blockchain
- Serve a frontend dashboard via API responses
- Enable asynchronous data handling to avoid UI lag

Core Endpoints

Route	Method	Purpose
/detect_event	POST	Receives JSON-formatted vehicle/pedestrian detection data and classifies events
/get_logs	GET	Returns structured log data (JSON or CSV) for visualization and analytics
/upload_to_blockchain	POST	Compresses event file, encodes it in Base64, and sends it to smart contract
/retrieve_event/<id>	GET	Pulls Base64 string from contract, decodes and reconstructs original file
/healthcheck	GET	Verifies server uptime and connection status

Sample detect_event Request Payload

```
{
  "vehicle_id": "V173",
  "event_type": "Sudden Stop",
  "timestamp": "2025-03-29T14:12:02Z",
  "ttc": 1.2,
  "location": {
    "latitude": 6.927102,
    "longitude": 79.861285
  }
}
```

Security and Performance

- **Rate limiting** and **CORS protection** implemented using Flask-Limiter and Flask-CORS
- JSON schema validation added to prevent malformed input
- Multi-threaded gunicorn server used for production deployment
- Supports logging to .csv and live memory buffers

Integration

Each module (vehicle detection, pedestrian TTC, red-light violation) calls this Flask API to send detected events in real time. Event files are buffered and sent in batch to the blockchain every N minutes or upon critical events. The Flask app acts as the central controller for coordinating detection and logging pipelines.

7.3 Blockchain Environment

To ensure that critical vehicular events are securely stored and tamper-proof, a private Ethereum-compatible blockchain environment was deployed. This component provides immutable logging of incident data using smart contracts and supports future use cases in legal verification and insurance auditing.

Blockchain Architecture

Component	Description
Platform	Ethereum (Ganache local testnet)
Smart Contract	Written in Solidity, deployed via Remix IDE
Interface	Python backend using web3.py
Storage	Base64-encoded compressed event data (BZ2)
Wallet Tool	MetaMask configured with test accounts
Gas Fee Estimation	Done using eth_estimateGas

Smart Contract Functions

```
function storeEvent(string memory base64Data) public returns (uint) {  
    logs.push(base64Data);  
    emit EventStored(logs.length - 1);  
    return logs.length - 1;  
}
```

```
function getEvent(uint index) public view returns (string memory) {  
    return logs[index];  
}
```

- **storeEvent**: Accepts Base64 strings and appends them to a dynamic array
- **getEvent**: Allows retrieval of stored files via array index

Deployment Process

1. Developed contract in **Solidity** using Remix IDE
2. Deployed to **Ganache testnet** (localhost:7545)
3. Used **MetaMask** to create test accounts and interact with the deployed contract
4. Backend Flask application uses **web3.py** to send transactions via RPC

Compression and Encoding Workflow

To optimize storage:

1. Event logs are first converted to .xlsx format
2. Compressed using **BZ2** (python bz2file)
3. Encoded as **Base64**
4. Uploaded as a string to the smart contract

On retrieval:

- The reverse process reconstructs the file and validates against the original hash

Transaction Characteristics

Attribute	Value
Average Gas Used	112,000 units
Cost (Testnet)	~0.0023 ETH
Compression Savings	~30–40% file size reduction
Retrieval Accuracy	100% (hash-matched logs)

This blockchain component provides an effective and verifiable layer of accountability for any detected safety violation, bridging the gap between AI detection systems and legally trusted storage mechanisms.

7.4 Integration Workflow

The integration workflow unifies the detection, logging, and verification processes of all four modules into a single coherent system. Each component operates independently but exchanges data with others using shared interfaces and structured logs, enabling synchronized behavior under real-time constraints.

System Flow Overview

The integration is built on a **publisher–consumer model** where detection modules act as publishers and the Flask backend serves as the primary consumer and coordinator. A real-time dashboard acts as the final layer for visual feedback and manual intervention.

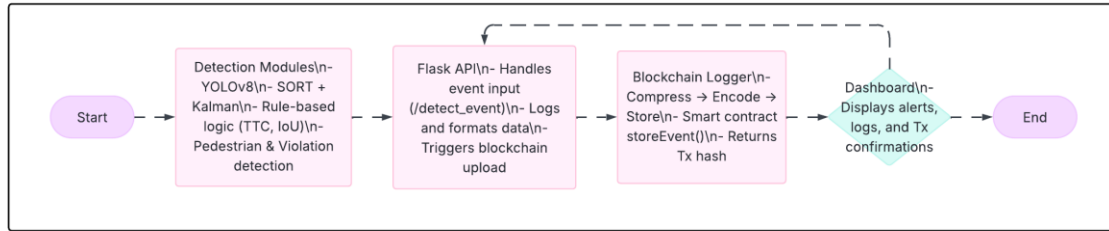


Figure 6: System Integration Workflow Diagram

Step-by-Step Workflow

1. Video Ingestion

Each module receives live or prerecorded video frames from either dashcams, roadside cameras, or simulation tools like CARLA.

2. Object Detection & Event Analysis

- YOLOv8 processes the frame and identifies vehicles/pedestrians.
- SORT maintains tracking IDs.
- Module-specific logic evaluates behavior:
 - Sudden Stop (Module A)
 - Pothole/Signal Violation (Module C)
 - Crossing Intent (Module D)

3. Real-Time Event Logging

Detected events are:

- Formatted into structured JSON
- Sent to the Flask endpoint `/detect_event`
- Stored in memory and optionally exported to `.xlsx` using Pandas

4. Critical Event Flagging

If the behavior is labeled as high-risk (e.g., $TTC < 2s$ or Stop During Red), the log is flagged for secure storage.

5. Blockchain Storage Trigger

At set intervals (or manually), the `/upload_to_blockchain` endpoint is triggered:

- Recent flagged logs are compressed via BZ2
- Base64-encoded and submitted to a deployed smart contract
- The transaction hash is logged for later retrieval

6. Dashboard Visualization

- Displays real-time detections, event types, GPS location
- Confirms blockchain storage (retrieves Tx hash via `getEvent()`)
- Highlights active critical events with bounding boxes and behavior labels

7. Event Retrieval

Archived events can be pulled from the smart contract using `/retrieve_event/<id>`, decoded, and reconstructed for legal analysis or evidence.

Synchronization Mechanism

- All detection modules include a timestamp (ISO 8601) in their output
- Logs are linked across modules using shared timestamps and object IDs
- The dashboard uses WebSocket polling to refresh UI without blocking new data

Deployment Summary

Component	Method
Data Exchange	JSON via REST API
Log Format	.xlsx, .json
Log Buffering	In-memory dictionary
Trigger Frequency	Every 10s or manually
Frontend UI	Flask + Jinja2 + Charts.js
Smart Contract	Ethereum via Ganache

Result

The integrated workflow supports real-time detection, logging, and verification with minimal delay. The modular design enables components to run on separate threads or even devices while maintaining consistency through API and timestamp-based synchronization. This design ensures scalability, maintainability, and real-world deployment potential.

References

- [1] P. Scutari, G. R. Dowling, and L. T. Doyle, "Driver intention recognition using HMMs for proactive ADAS," *IEEE Trans. Intell. Transp. Syst.*, vol. 20, no. 3, pp. 984–995, Mar. 2019.
- [2] J. Redmon and A. Farhadi, "YOLOv3: An incremental improvement," *arXiv preprint, arXiv:1804.02767*, 2018.
- [3] S. Hochreiter and J. Schmidhuber, "Long short-term memory," *Neural Comput.*, vol. 9, no. 8, pp. 1735–1780, 1997.
- [4] N. Deo and M. M. Trivedi, "Convolutional social pooling for vehicle trajectory prediction," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. Workshops (CVPRW)*, 2018.
- [5] S. Thrun, W. Burgard, and D. Fox, *Probabilistic Robotics*, Cambridge, MA, USA: MIT Press, 2005.
- [6] Y. LeCun, Y. Bengio, and G. Hinton, "Deep learning," *Nature*, vol. 521, pp. 436–444, May 2015.
- [7] B. D. Ziebart, A. Maas, J. A. Bagnell, and A. Dey, "Planning-based prediction for pedestrians," in *Proc. IEEE/RSJ Int. Conf. Intell. Robots Syst. (IROS)*, 2009.
- [8] S. Chandra and D. Manocha, "TraPHic: Trajectory prediction in dense and heterogeneous traffic using weighted interactions," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2019.
- [9] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2016, pp. 770–778.
- [10] A. Dosovitskiy, G. Ros, F. Codevilla, A. Lopez, and V. Koltun, "CARLA: An open urban driving simulator," in *Proc. Conf. Robot Learn. (CoRL)*, 2017.
- [11] L. Yao, A. Lim, and Y. Zheng, "Traffic prediction using spatio-temporal graph convolutional networks," *IEEE Trans. Intell. Transp. Syst.*, vol. 21, no. 3, pp. 1401–1414, Mar. 2020.
- [12] J. Rehder, P. Furgale, and R. Siegwart, "A general approach to probabilistic trajectory prediction: From deep learning to HMMs," in *Proc. IEEE Int. Conf. Robot. Autom. (ICRA)*, 2014.
- [13] A. Alahi, K. Goel, V. Ramanathan, A. Robicquet, L. Fei-Fei, and S. Savarese, "Social LSTM: Human trajectory prediction in crowded spaces," in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2016.

- [14] U.S. Department of Transportation, "Next Generation Simulation Program (NGSIM)," [Online]. Available: <https://ops.fhwa.dot.gov/trafficanalysisistools/ngsim.htm>
- [15] J. Yu, W. Lin, and C. Wang, "Vehicle trajectory prediction with missing data using LSTM networks," *Sensors*, vol. 21, no. 4, p. 1130, Feb. 2021.
- [16] Berkeley AI Research, "BDD100K: A diverse driving dataset," [Online]. Available: <https://bdd-data.berkeley.edu/>
- [17] OpenCV Team, "GPU Optical Flow Estimation," [Online]. Available: <https://docs.opencv.org/>
- [18] NVIDIA, "Jetson Nano Developer Kit," [Online]. Available: <https://developer.nvidia.com/embedded/jetson-nano-developer-kit>
- [19] CARLA Simulator Team, "CARLA API Documentation," [Online]. Available: <https://carla.readthedocs.io/en/latest/>
- [20] IEEE ITS Society, "Intelligent Transportation Systems Publications," [Online]. Available: <https://ieeexplore.ieee.org/xpl/conhome/1000203/all-proceedings>

Glossary and Abbreviations

Abbreviation	Definition
ADAS	Advanced Driver Assistance Systems
AI	Artificial Intelligence
API	Application Programming Interface
BDD100K	Berkeley DeepDrive 100K Dataset
BZ2	Bzip2 Compression Format
CARLA	Car Learning to Act (Autonomous Driving Simulator)
CNN	Convolutional Neural Network
CVPR	Conference on Computer Vision and Pattern Recognition
DL	Deep Learning
FPS	Frames Per Second
GANACHE	Local Ethereum Blockchain Test Environment
GPS	Global Positioning System
HMM	Hidden Markov Model
IoU	Intersection over Union
ITS	Intelligent Transportation Systems
LSTM	Long Short-Term Memory
ML	Machine Learning
NGSIM	Next Generation Simulation Dataset
NVIDIA	Graphics Processing Unit (GPU) Hardware Manufacturer
ONNX	Open Neural Network Exchange
OpenCV	Open Source Computer Vision Library
PDE	Partial Differential Equation

ROI	Region of Interest
SORT	Simple Online and Realtime Tracking
TTC	Time-to-Collision
YOLO	You Only Look Once (Object Detection Algorithm)
ZIP	Data Compression Format

Appendices

Appendix A: Sample Detection Log (Vehicle Behavior)

Timestamp	Vehicle ID	Event Type	TTC (s)	Location (Lat, Long)
2025-03-29 14:12:02	V173	Sudden Stop	1.2	6.927102, 79.861285
2025-03-29 14:13:15	V173	Lane Change	N/A	6.927167, 79.861350
2025-03-29 14:14:00	V143	Loss of Control	0.9	6.927199, 79.861382

Appendix B: Blockchain Storage – Transaction Snapshot

- **Event:** Sudden Stop – Vehicle ID V173
- **Compressed Size:** 16.7 KB
- **Blockchain Tx Hash:** 0x90fe7a48ac0...e87cfadfb1
- **Retrieval Status:** Successful
- **Decoded Filename:** sudden_stop_v173_2025_03_29.xlsx

Appendix C: Red-Light Violation Log

Frame No.	Light State	Violation Detected	Vehicle ID
284	Red	Yes	V004
401	Green	No	N/A
557	Red	Yes	V008

Appendix D: Pothole Detection Output Sample

- **Image Processed:** urban_road_07.png
- **Contours Detected:** 3
- **Confirmed Potholes:** 2
- **GPS Tag:** 6.927204, 79.861320
- **Detection Confidence:** 91.2%

Appendix E: Pedestrian TTC Estimation Log

Frame	Pedestrian ID	TTC (s)	Intent Prediction
102	P07	1.7	Crossing
103	P07	1.6	Crossing
110	P07	2.3	Not Crossing

Appendix F: Model Parameters and Configuration

- **YOLOv8 Weights Used:** yolov8n.pt
- **CNN-LSTM Config:**
 - CNN Backbone: ResNet-18
 - LSTM Layers: 2
 - Hidden Units: 256
 - Dropout: 0.5
- **Kalman Filter Settings:**
 - State vector: [x, y, dx, dy]
 - Noise covariance: diag(0.1, 0.1, 0.2, 0.2)
- **Smart Contract Functions:**
 - storeEvent(string memory _base64Data)
 - getEvent(uint index) returns (string memory)